

ECG Classification Using Neural Networks

Sai Yadavalli

Content

1 Problem



2 Data



3 Algorithm



4 Training




5 Results



6 Future



Problem



- No way to automatically classify different types of heartbeats from ECG signals
- End Goal: Early, accurate detection of arrhythmias, reducing the burden on healthcare professionals by providing automated diagnostic support.

Data Description

109446

rows/samples (training +
validation)

187

Features: 187 values per heartbeat
(ECG signal over time)

5

Labels: Normal → 0, Supraventricular
ectopic → 1, Ventricular ectopic → 2,
Fusion → 3, Unknown → 4

Source

MIT-BIH

Arrhythmia
Dataset

PhysioNet

Kaggle

Each sample is a 1D vector
representing a fixed-length ECG
waveform segment.

[Link to Kaggle Data Set](#)

Convolutional Neural Network (CNN) Model

1



Libraries

- Implemented from scratch in NumPy—no deep learning libraries used
- Pandas for data processing
- Scikit-Learn for metrics

2



Implementation

- 1D Convolution layer extracts local ECG patterns
- Max pooling reduces feature dimensionality
- Softmax layer for probability output
- ReLU activation and backpropagation—manually implemented

3



Architecture (Fully Connected)

Conv → Pool → Dense → Output

Training & Testing

Data was split into...

- 80% Training
- 20% Testing/Validation
- Not Standardized



Manual implementation of...

- Forward and backward passes
- Stochastic gradient descent
- Loss: Cross-Entropy

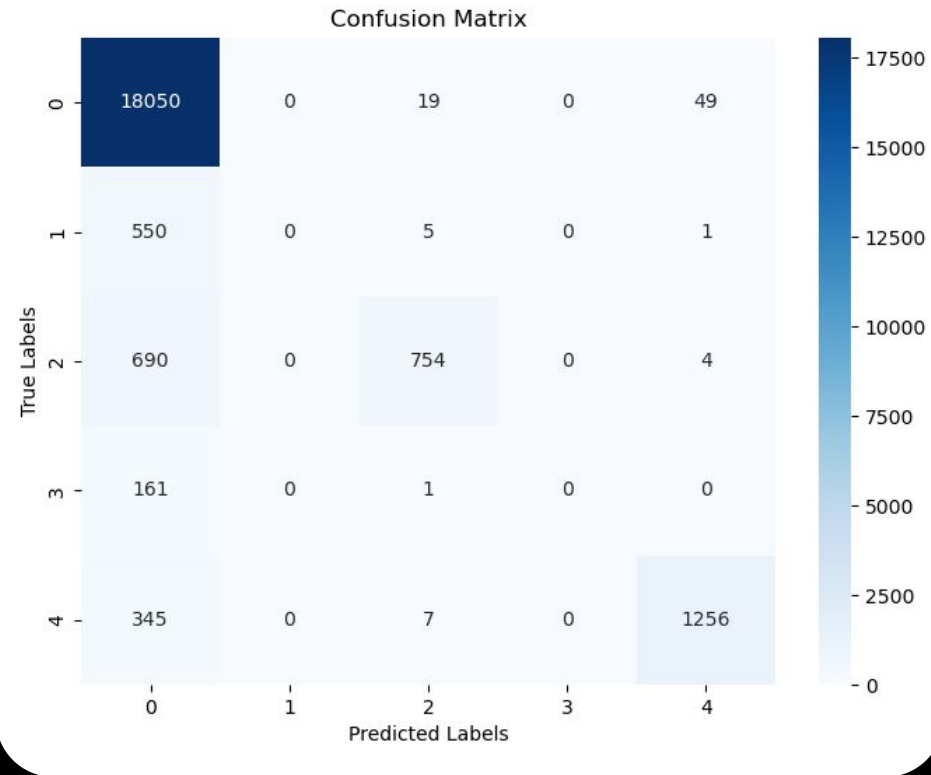
$$H(p, q) = - \sum_{x \in \mathcal{X}} p(x) \log q(x).$$



Let's work together

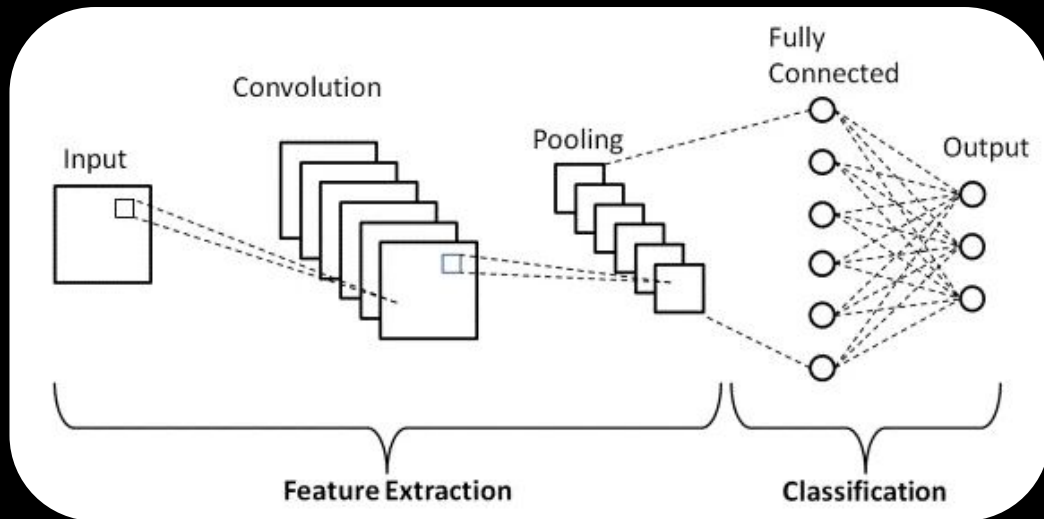
Results & Evaluation

- Evaluated on test set using classification metrics:
 - Accuracy: 91.63%
 - Precision: 88.85%
 - Recall: 91.63%
 - F1 Score: 89.59%



What Did I Learn?

- Even simple CNNs can achieve strong performance on time-series data in Arrhythmia Classification
- Implementing backpropagation, training logic, and loss function deepened understanding of ML fundamentals
- Future improvements:
 - Add more layers
 - Adjust Parameters
 - `conv_filters=6,`
`kernel_size=5,`
`pool_size=2,`
`hidden_units=32,`
`learning_rate=0.01,`
`epochs=10,`
`batch_size=32`
 - Try Scaling



Thank you

Any
Questions?